

XV-6 : Traps

Traps Processing (hard ware)

Traps

- Exception
 - ↳ sys call
 - ↳ program error.
- Interrupt.

Example:

- System call
"ecall" instruction

- Program Error.
 - Illegal instruction,
 - Alignment errors
- Device Interrupt.
 - interrupt can occur when we are executing in user mode or supervisor mode

After trap happens, regardless of any mode, we begin executing some handler code running in supervisor mode.



it is a CSR

contains the address of the "Handler" code.

There are two pieces of code in xv6 that handle traps.

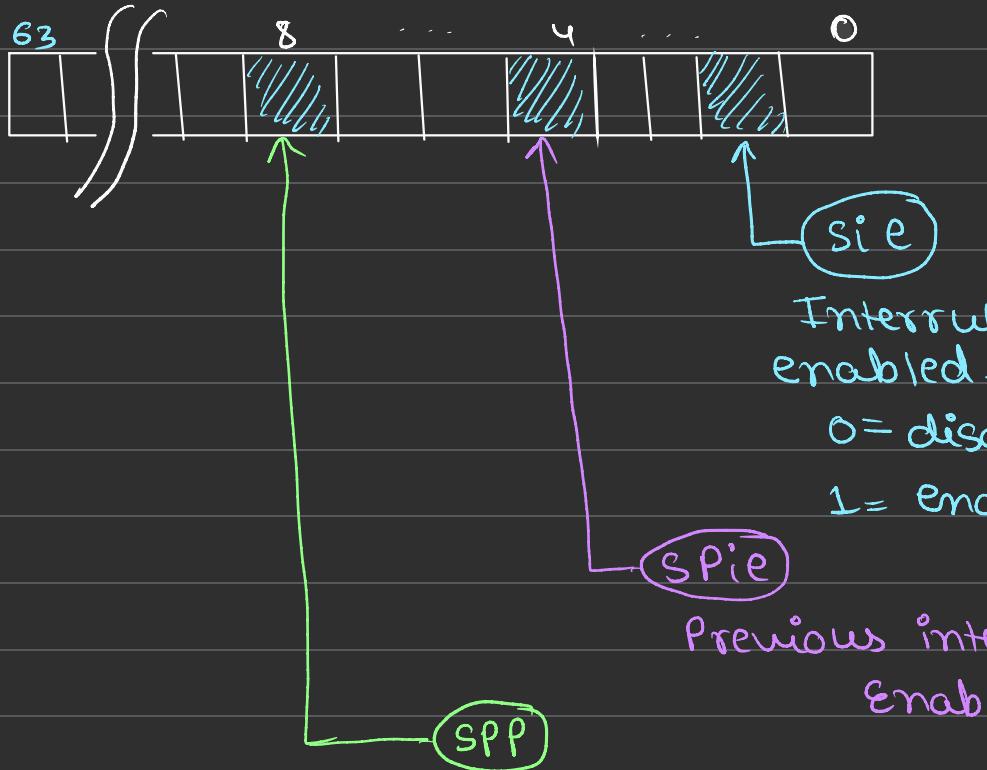
- kernelvec

Handles traps during code in kernel mode.

- uservec

Handles traps that happens when running in user mode.

sstatus



Interruptions enabled.

0 = disabled

1 = enabled

SPP

Previous privilege level
(Previous mode)

0 → user mode

1 → supervisor mode

What happens when a Trap occurs?

Are interrupts disabled?

→ The interrupt remains "Pending" until interrupts are re-enabled.

↳ If it is an exception it's going to be handled immediately regardless of whether interrupts are disabled or not.

when traps occurs, The hardware does this.

SEPC \leftarrow PC

PC \leftarrow stvec

scause \leftarrow ...

stval \leftarrow ... additional info. . .

//hardware will save previous mode
Sstatus.SPP \leftarrow Previous mode

Sstatus.SPIE \leftarrow Sstatus.SIE

↑
Save previous "interrupts enabled bit."

Sstatus.SIE \leftarrow \emptyset

Disable Interrupt

// change the mode to supervisor
mode ← supervisor. (if its not supervisor)

And then the first instruction of the trap handler executes.

Handler code do something.

Later, the handler returns to the interrupted code.

sret

sstatus.sie ← sstatus.spi.e

mode ← sstatus.spp.

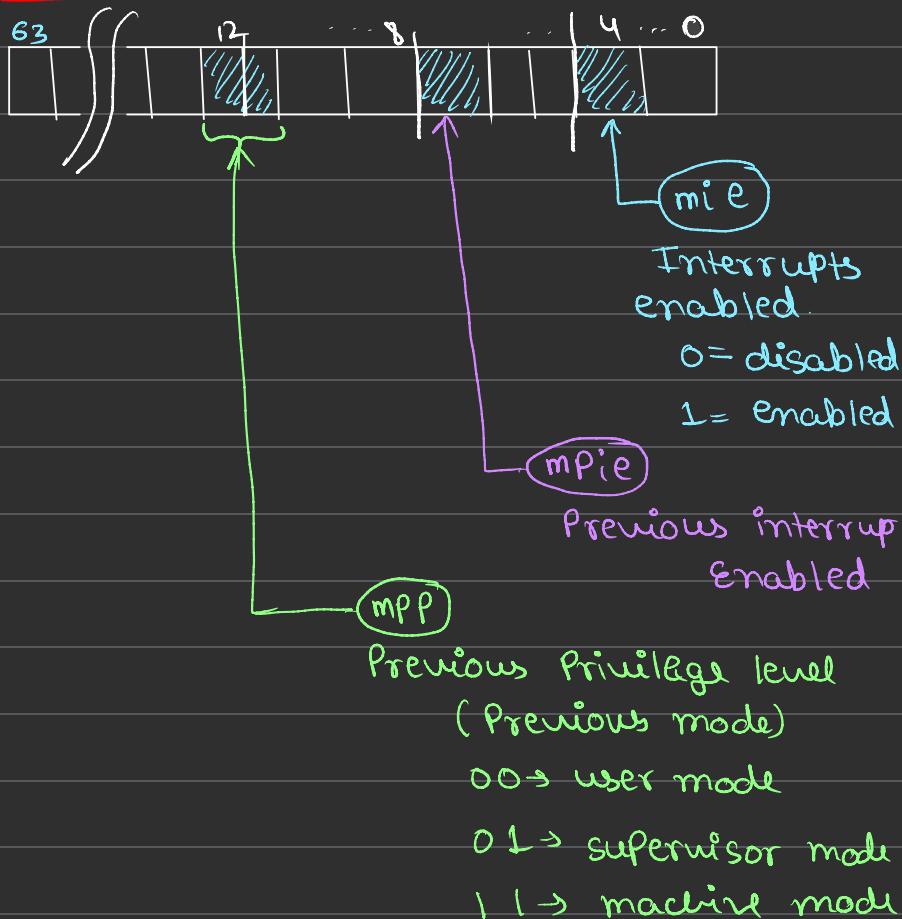
PC ← sepc

That's what happens in supervisor mode and most of the runs in the supervisor mode

But there's a little bit of the kernel that runs in machine mode.

Trap Processing in machine mode

mstatus



- ① Only Interrupt we deal with at machine mode level is Timer Interrupt.
[All others go to ("delegated") supervisor mode.]
- ② Interrupts are always enabled at machine mode.
(whenever a timer interrupt occurs it will be handled, whenever any trap of other kind happens it will be automatically delegated to supervisor mode)
- ③ On timer Interrupt a machine mode handler will run and that code:
 - force a "software" Interrupt to supervisor mode
 - re-enable Interrupts.

- Return To ... the interrupted code ...

Trap Processing in machine mode.

{Same Idea}

only cause : Timer Interrupts.

All other types are "Delegated" to supervisor mode

mtvec: CSR containing address of handler.

→ in xv6, that code is timervec function

Hardware Actions

mePC \leftarrow PC

PC \leftarrow mtvec

mcause \leftarrow ... } ignored.
mtval \leftarrow ... }

$mstatus.mPP \leftarrow$ Previous mode

$mstatus.MIE \leftarrow mstatus.MIE$

$mstatus.MIE \leftarrow \phi$

mode \leftarrow machine

The handler will handle the interrupt.

The handler will trigger / cause an interrupt at the supervisor level.

m ret

$mstatus.MIE \leftarrow mstatus.MIE$

mode $\leftarrow mstatus.mPP$

PC $\leftarrow mepc$.

Next instruction

- Interrupts Disabled

"Software" interrupt pending

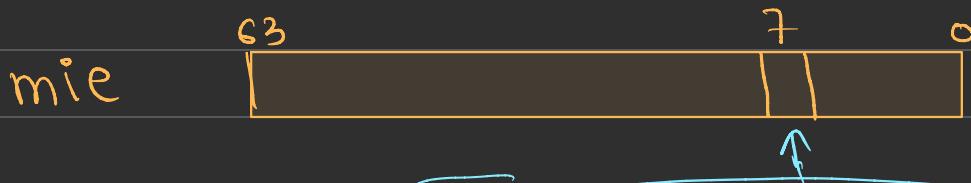
- Interrupts Enabled

A trap occurs at supervisor

level

Global "Interrupts enabled" bit in mstatus, sstatus CSR.

more selective control.

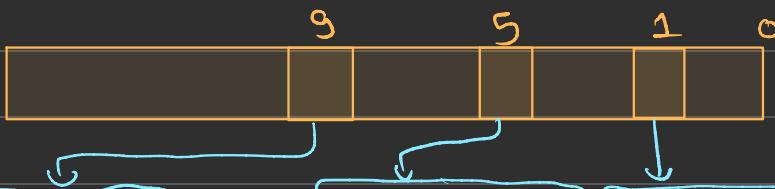


"MTIE"

- 1: Enable Timer interrupts.
- 0: Disable Timer interrupt

Set during machine mode initialization

sie



"SSIE"
Software interrupt

"STIE"
Timer interrupts

"SEIE"
Device Interrupts

SIP: same as above.

- 1: An interrupt is pending
- 0: No interrupt.

Trap Delegation

All traps in RISC-V architecture will go to machine mode handlers

But there is also a facility to sort of bypass the machine mode handler and go straight to handler that's executing at supervisor mode level.

XV6 delegates all traps to supervisor mode.

There are two registers:

Hard ware Delegation Registers

medeleg which exceptions to delegate
mideleg which interrupts to delegate
(CSRs)

XV6 delegates all traps to supervisor mode.

Timer Interrupt cannot be delegated. It is handled differently.

medeleg ,mideleg are initialized during set up (in "start" in machine mode).

when a trap occurs

- No machine mode code.
- See Trap Processing for supervisor mode

{ Trap }

{ SRET }

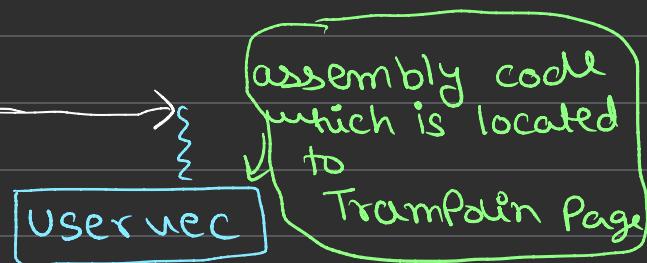
Usercode

Road MAP

- So this trampoline page recall is mapped into all address spaces both user & kernel

Trap
init ← disabled
model ← S
sepc ← PC
scawse ← ...

- Timer Interrupt
- Device Interrupt.
- syscall
- exception(error)



Saves regs,PC in trapframe
Restore Kernel's SP,TP.
satp ← Kernel Page table
Jump to... {SP → Stack Pointer
every threads needs its own stack, so we have to initialize the stack Pts. }
{TP → we also initialize TP which contains core number }
{satp : Pointer to the Kernel Page table}

usertrap()

stevec ← Kernel's Trap vector.

we look at the cause reg.

Exception

Print exit()

Device

deviceintr()

{ if killed was set exit() }

Syscall

{ inter-Enabled
syscall() }

{ if killed was set ... exit() }

Timer

yield()



usertrapret()

{
 imts ← Disabled
 stvec ← uservec
 Save SP, tP.
 sepc ← save PC

Userret

{
 satp ← user's PageTable
 restore user's Regs.
 sstatus ← {
 SPP = "Umode"
 sret SP + E = "Enabled"

←
{}
SRET